

Supervised Learning Models for Customer Classification

```
In [5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score,
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
```

```
In [6]: cleaned_data = pd.read_csv("data/02_removed_outliers_redundant.csv", index_col=0)
clustered_data = pd.read_csv("results/clustering/clustered_data_results.csv", index
```

```
In [19]: clustered_data.head()
```

```
Out[19]:
```

	Education	Marital_Status	Income	Recency	Response	Age	Customer_Since	T
0	-0.938689	-1.351057	0.323276	0.309449	2.391652	1.020547	1.528805	
1	-0.938689	-1.351057	-0.252104	-0.382368	-0.418121	1.278260	-1.187852	
2	-0.938689	0.740161	0.980665	-0.797458	-0.418121	0.333312	-0.204916	
3	-0.938689	0.740161	-1.213087	-0.797458	-0.418121	-1.298871	-1.059428	
4	1.065316	0.740161	0.330838	1.554719	-0.418121	-1.041158	-0.950762	

```
In [ ]: # Prepare features and set cluster as target
X_all = clustered_data.drop('Cluster', axis=1)
y = clustered_data['Cluster']
```

```
In [8]: # define feature sets as mentioned in the project proposal
# all features (for existing customers with purchase history)
all_features = X_all.columns.tolist()
```

```
In [9]: # Remove purchase-related and spending-related features
purchase_related = ['Total_Spent', 'RatioWines', 'RatioFruits', 'RatioMeatProducts',
                    'RatioFishProducts', 'RatioSweetProducts', 'RatioGoldProds',
                    'Total_Accepted_Campaign', 'Total_Purchase', 'Total_Web_Engageme
```

```
In [10]: # separate between demo and other type of features
demographic_features = [f for f in all_features if f not in purchase_related]

print("Feature sets:")
print(f"All features ({len(all_features)}): {all_features}")
print(f"Demographic features ({len(demographic_features)}): {demographic_features}")
```

Feature sets:

All features (17): ['Education', 'Marital_Status', 'Income', 'Recency', 'Response', 'Age', 'Customer_Since', 'Total_Spent', 'RatioWines', 'RatioFruits', 'RatioMeatProducts', 'RatioFishProducts', 'RatioSweetProducts', 'RatioGoldProds', 'Total_Accepted_Campaign', 'Total_Purchase', 'Total_Web_Engagement']

Demographic features (7): ['Education', 'Marital_Status', 'Income', 'Recency', 'Response', 'Age', 'Customer_Since']

```
In [29]: # Define models to test
models = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'Random Forest': RandomForestClassifier(random_state=42, n_estimators=100),
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),
    'SVM': SVC(random_state=42, probability=True),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Neural Network': MLPClassifier(random_state=42, max_iter=1000)
}
```

```
In [ ]: # evaluate all 6 models
def evaluate_models(X, y, feature_set_name):
    print(f"Evaluating Models with {feature_set_name}")

    # split the data 80: 20
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                         random_state=42, stratify=y)

    results = {}

    for name, model in models.items():
        print(f"\nTraining {name}...")

        # cross-validation 5 fold
        cv_scores = cross_val_score(model, X_train, y_train, cv=5, scoring='accuracy')

        # fit and predict
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        y_pred_proba = model.predict_proba(X_test)[:, 1]

        # Calculate metrics
        test_accuracy = model.score(X_test, y_test)
        roc_auc = roc_auc_score(y_test, y_pred_proba)

        results[name] = {
            'model': model,
            # 'scaler': scaler,
            'cv_mean': cv_scores.mean(),
            'cv_std': cv_scores.std(),
```

```
        'test_accuracy': test_accuracy,  
        'roc_auc': roc_auc,  
        'y_test': y_test,  
        'y_pred': y_pred,  
        'y_pred_proba': y_pred_proba,  
        'classification_report': classification_report(y_test, y_pred)  
    }  
  
    print(f"   CV Accuracy: {cv_scores.mean():.3f} (+/- {cv_scores.std() * 2:.3f}  
    print(f"   Test Accuracy: {test_accuracy:.3f}")  
    print(f"   ROC AUC: {roc_auc:.3f}")  
  
    return results, X_train, X_test, y_train, y_test
```

```
In [22]: # Evaluate models with all features  
results_all = evaluate_models(X_all, y, "All Features")[0]  
  
# Evaluate models with demographic features only  
X_demo = X_all[demographic_features]  
results_demo = evaluate_models(X_demo, y, "Demographic Features Only")[0]
```

Evaluating Models with All Features

Training Logistic Regression...

CV Accuracy: 0.993 (+/- 0.008)

Test Accuracy: 0.998

ROC AUC: 1.000

Training Random Forest...

CV Accuracy: 0.963 (+/- 0.015)

Test Accuracy: 0.954

ROC AUC: 0.995

Training Gradient Boosting...

CV Accuracy: 0.967 (+/- 0.013)

Test Accuracy: 0.961

ROC AUC: 0.994

Training SVM...

CV Accuracy: 0.979 (+/- 0.010)

Test Accuracy: 0.975

ROC AUC: 0.999

Training K-Nearest Neighbors...

CV Accuracy: 0.948 (+/- 0.022)

Test Accuracy: 0.934

ROC AUC: 0.983

Training Neural Network...

CV Accuracy: 0.994 (+/- 0.008)

Test Accuracy: 0.991

ROC AUC: 1.000

Evaluating Models with Demographic Features Only

Training Logistic Regression...

CV Accuracy: 0.841 (+/- 0.047)

Test Accuracy: 0.828

ROC AUC: 0.891

Training Random Forest...

CV Accuracy: 0.857 (+/- 0.015)

Test Accuracy: 0.856

ROC AUC: 0.905

Training Gradient Boosting...

CV Accuracy: 0.853 (+/- 0.037)

Test Accuracy: 0.826

ROC AUC: 0.897

Training SVM...

CV Accuracy: 0.845 (+/- 0.053)

Test Accuracy: 0.838

ROC AUC: 0.897

Training K-Nearest Neighbors...

CV Accuracy: 0.830 (+/- 0.037)

Test Accuracy: 0.824

ROC AUC: 0.856

Training Neural Network...

CV Accuracy: 0.838 (+/- 0.044)

Test Accuracy: 0.842

ROC AUC: 0.895

```
In [23]: # compare all models
comparison_data = []

for feature_set, results in [("All Features", results_all), ("Demographic Only", re
    for model_name, result in results.items():
        comparison_data.append({
            'Feature Set': feature_set,
            'Model': model_name,
            'CV Accuracy': result['cv_mean'],
            'Test Accuracy': result['test_accuracy'],
            'ROC AUC': result['roc_auc']
        })

comparison_df = pd.DataFrame(comparison_data)
print(comparison_df.round(3))
```

	Feature Set	Model	CV Accuracy	Test Accuracy	ROC AUC
0	All Features	Logistic Regression	0.993	0.998	1.000
1	All Features	Random Forest	0.963	0.954	0.995
2	All Features	Gradient Boosting	0.967	0.961	0.994
3	All Features	SVM	0.979	0.975	0.999
4	All Features	K-Nearest Neighbors	0.948	0.934	0.983
5	All Features	Neural Network	0.994	0.991	1.000
6	Demographic Only	Logistic Regression	0.841	0.828	0.891
7	Demographic Only	Random Forest	0.857	0.856	0.905
8	Demographic Only	Gradient Boosting	0.853	0.826	0.897
9	Demographic Only	SVM	0.845	0.838	0.897
10	Demographic Only	K-Nearest Neighbors	0.830	0.824	0.856
11	Demographic Only	Neural Network	0.838	0.842	0.895

All features outperform demographic only.

```
In [24]: # Create visualization of model performance
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

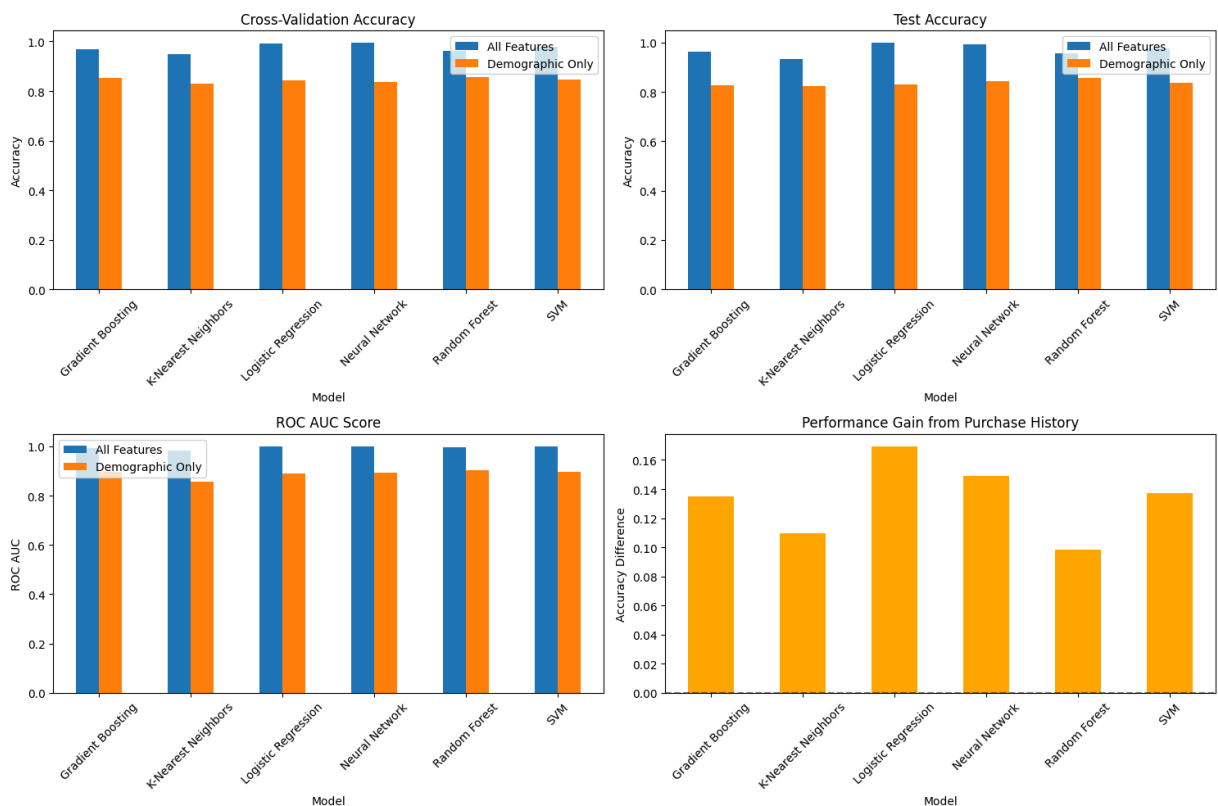
# CV Accuracy comparison
pivot_cv = comparison_df.pivot(index='Model', columns='Feature Set', values='CV Acc
pivot_cv.plot(kind='bar', ax=axes[0, 0], rot=45)
axes[0, 0].set_title('Cross-Validation Accuracy')
axes[0, 0].set_ylabel('Accuracy')
axes[0, 0].legend()

# Test Accuracy comparison
pivot_test = comparison_df.pivot(index='Model', columns='Feature Set', values='Test
pivot_test.plot(kind='bar', ax=axes[0, 1], rot=45)
axes[0, 1].set_title('Test Accuracy')
axes[0, 1].set_ylabel('Accuracy')
axes[0, 1].legend()
```

```
# ROC AUC comparison
pivot_roc = comparison_df.pivot(index='Model', columns='Feature Set', values='ROC AUC')
pivot_roc.plot(kind='bar', ax=axes[1, 0], rot=45)
axes[1, 0].set_title('ROC AUC Score')
axes[1, 0].set_ylabel('ROC AUC')
axes[1, 0].legend()

# Performance difference (ALL Features - Demographic Only)
performance_diff = pivot_test['All Features'] - pivot_test['Demographic Only']
performance_diff.plot(kind='bar', ax=axes[1, 1], color='orange', rot=45)
axes[1, 1].set_title('Performance Gain from Purchase History')
axes[1, 1].set_ylabel('Accuracy Difference')
axes[1, 1].axhline(y=0, color='black', linestyle='--', alpha=0.5)

plt.tight_layout()
plt.savefig('plots/model_comparison.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
In [25]: # Find best models for each feature set
best_all = max(results_all.items(), key=lambda x: x[1]['test_accuracy'])
best_demo = max(results_demo.items(), key=lambda x: x[1]['test_accuracy'])

print(f"\nBest model with all features: {best_all[0]} (Accuracy: {best_all[1]['test_accuracy']})")
print(f"Best model with demographic features: {best_demo[0]} (Accuracy: {best_demo[1]['test_accuracy']})")

# stat for the best models
print(f"\nBest All Features Model: {best_all[0]}")
print(best_all[1]['classification_report'])

print(f"\nBest Demographic Model: {best_demo[0]}")
print(best_demo[1]['classification_report'])
```

Best model with all features: Logistic Regression (Accuracy: 0.998)
 Best model with demographic features: Random Forest (Accuracy: 0.856)

Best All Features Model: Logistic Regression

	precision	recall	f1-score	support
0	1.00	1.00	1.00	275
1	1.00	0.99	1.00	162
accuracy			1.00	437
macro avg	1.00	1.00	1.00	437
weighted avg	1.00	1.00	1.00	437

Best Demographic Model: Random Forest

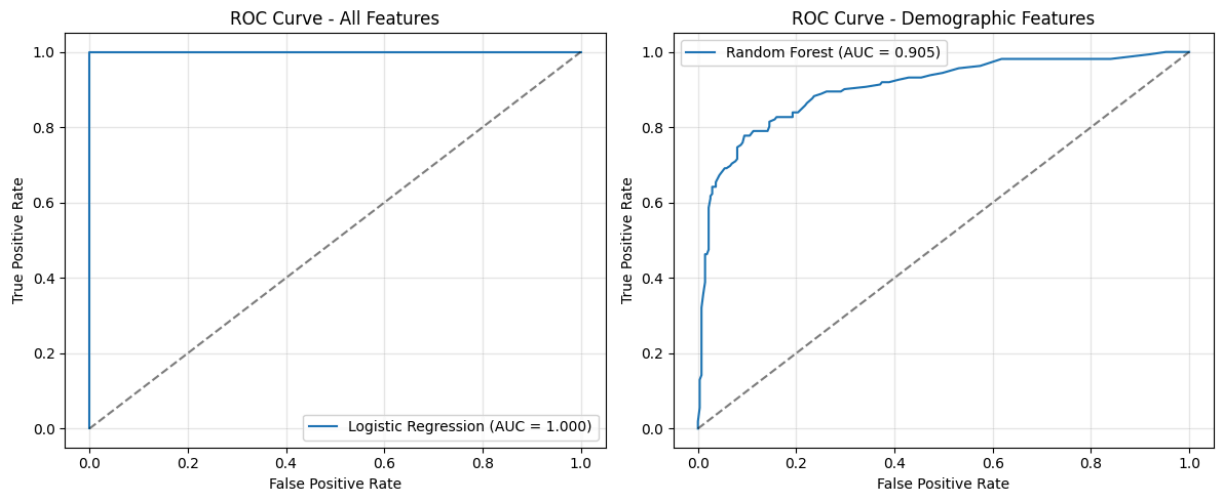
	precision	recall	f1-score	support
0	0.86	0.92	0.89	275
1	0.85	0.75	0.79	162
accuracy			0.86	437
macro avg	0.85	0.83	0.84	437
weighted avg	0.86	0.86	0.85	437

```
In [26]: # ROC Curves for best models
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
fpr, tpr, _ = roc_curve(best_all[1]['y_test'], best_all[1]['y_pred_proba'])
plt.plot(fpr, tpr, label=f'{best_all[0]} (AUC = {best_all[1]["roc_auc"]:.3f})')
plt.plot([0, 1], [0, 1], 'k--', alpha=0.5)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - All Features')
plt.legend()
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
fpr, tpr, _ = roc_curve(best_demo[1]['y_test'], best_demo[1]['y_pred_proba'])
plt.plot(fpr, tpr, label=f'{best_demo[0]} (AUC = {best_demo[1]["roc_auc"]:.3f})')
plt.plot([0, 1], [0, 1], 'k--', alpha=0.5)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Demographic Features')
plt.legend()
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('plots/best_models_roc_curves.png', dpi=300, bbox_inches='tight')
plt.show()
```




```
In [27]: # confusion matrices
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

cm_all = confusion_matrix(best_all[1]['y_test'], best_all[1]['y_pred'])
sns.heatmap(cm_all, annot=True, fmt='d', ax=axes[0], cmap='Blues')
axes[0].set_title(f'Confusion Matrix - {best_all[0]} (All Features)')
axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('Actual')

cm_demo = confusion_matrix(best_demo[1]['y_test'], best_demo[1]['y_pred'])
sns.heatmap(cm_demo, annot=True, fmt='d', ax=axes[1], cmap='Blues')
axes[1].set_title(f'Confusion Matrix - {best_demo[0]} (Demographic)')
axes[1].set_xlabel('Predicted')
axes[1].set_ylabel('Actual')

plt.tight_layout()
plt.savefig('plots/best_models_confusion_matrices.png', dpi=300, bbox_inches='tight')
plt.show()
```



Both models achieved high accuracy, especially with full features (up to 99.8% with logistic regression), validating the segmentation. Models can effectively classify customers into the two segments. Purchase history significantly improves classification accuracy.. Using only

demographic data still yields reasonable performance, supporting targeted marketing strategies.